

GRETA WARREN

Postdoctoral Researcher, Department of Computer Science, University of Copenhagen

grwa@di.ku.dk ♦ gretawarren.github.io ♦ [ORCID: 0000-0002-3804-2287](https://orcid.org/0000-0002-3804-2287)

Research interests: Human-AI interaction, Explainable AI, Natural Language Processing, Fact-checking

EMPLOYMENT

University of Copenhagen Copenhagen, Denmark
Postdoctoral Researcher, Department of Computer Science Mar 2024 - Present

- Conducting interdisciplinary HCI and NLP research using qualitative and quantitative methods to understand and address diverse stakeholder needs for explainable AI
- Advisors: Prof. Isabelle Augenstein and Prof. Irina Shklovski

Trinity College Dublin Dublin, Ireland
Research Assistant, Trinity College Institute of Neuroscience Sep 2018 - Aug 2019

- Recruited participants and collected electroencephalography (EEG), pupillometric and neuropsychological data, conducted data entry, processing and analysis, compiled hospital research ethics applications, delivered talk on brain injury research to patients and staff in National Rehabilitation Hospital, Dublin
- Advisor: Prof. Paul Dockree

EDUCATION

University College Dublin Dublin, Ireland
Ph.D., School of Computer Science Sep 2019 - Jan 2024

- Thesis: *User-centred counterfactual explanations for explainable AI*
- Supervisors: Prof. Mark Keane and Prof. Ruth Byrne.
- Examiners: Prof. Tim Miller (University of Queensland) and Prof. David Coyle (University College Dublin)
- Awarded €6000+ in travel scholarships (ACM SIGCHI, ACM FAccT, ICCBR)
- Collaborated with Accenture Labs, The Dock, Dublin on applied explainable AI projects (2022 - 2023)

Trinity College Dublin Dublin, Ireland
B.A. (Hons), Psychology Sep 2014 - May 2018

- Thesis: *Effects of additive and subtractive counterfactual thinking on choice of political candidate*
- Supervisor: Prof. Ruth Byrne

TEACHING EXPERIENCE

Thesis Advisor Jan 2025-
Department of Computer Science, University of Copenhagen

- Stefan Gottlieb Kramer, MSc. in IT and Cognition. Primary supervisor for Thesis: *Using Graph Neural Networks for Usable, Explainable Source Reliability Determination in Fact-checking*. Jan 2025-Sep 2025. Grade: 10/12

Internal Censor/Examiner Jun 2025-
Department of Computer Science, University of Copenhagen

- Assessed and graded MSc. student thesis projects in Department of Computer Science.

Teaching Assistant Sep 2020 - Dec 2022
School of Computer Science, University College Dublin

- 2020-2022 Text Analytics, MSc course, coordinated demonstrators and administration, prepared lab materials, facilitated lab sessions and graded assignments
- 2021-2022 Data Science in Python, MSc course, prepared lab materials and facilitated lab sessions
- 2020-2022 Introduction to Programming, BSc course, prepared lab materials, facilitated lab sessions and graded assignments
- 2020-2021 Programming I & Object-Oriented Programming, MSc course, prepared lab materials, facilitated lab sessions and graded assignments

HONOURS & AWARDS

DAAD AInet Fellowship , German Federal Ministry of Research	Nov 2025
Best Paper Award , Pre-ACL Workshop, Danish Data Science Academy	July 2025
Gary Marsden Travel Award , ACM SIGCHI	€4,000 Mar 2023
Doctoral Consortium Scholarship , ICCBR 2022	€500 Sep 2022
Travel Award , ACM FAccT 2022	€1,500 Jun 2022
1st Prize, Arthur Cox Alternative Perspectives Essay Competition , Trinity College Law Review	May 2018
First Class Award , Trinity College Dublin	Sep 2015
Entrance Exhibition , Trinity College Dublin	€150 Sep 2014
All Ireland Scholarship , All Ireland Scholarship Trust	€27,000 Sep 2014

PROFESSIONAL SERVICE AND VOLUNTEERING

Associate Chair

ACM CHI Conference on Human Factors in Computing Systems <i>Understanding People: Statistical & Quantitative Methods Subcommittee</i>	2026
Nordic Conference on Human-Computer Interaction (NordiCHI)	2026

Grant Proposal Reviewer

French National Research Agency (ANR), Generic Call for Proposals (AAPG) Artificial Intelligence and Data Science panel	2026
--	------

Invited Contributor

Meta Oversight Board Roundtable (December 2025): Assessing Meta's plans to expand Community Notes. Work cited in Oversight Board's [advisory opinion](#) for Meta

Programme Committee Member / Reviewer

Conferences: ACM Conference on Fairness, Accountability, and Transparency (**FAccT**) (2026-), ACM **CHI** Conference on Human Factors in Computing Systems* (2025-) Association for Computational Linguistics (ACL) Rolling Review (**ARR**)* (2025-), Annual Meeting of the Cognitive Science Society (**CogSci**) (2024-), ACM Conference on Intelligent User Interfaces (**IUI**)* (2023-), AAAI/ACM Conference on Artificial Intelligence, Ethics, and Society (**AIES**) (2022-), Irish Conference on Artificial Intelligence and Cognitive Science (AICS) (2021)

Journals: Harvard Kennedy School Misinformation Review (2026), International Journal of Human-Computer Interaction (2025), Philosophy & Technology (2024), Decision Support Systems (2023), Expert Systems (2022)

Workshops and Symposia: European Symposium on Artificial Neural Networks, Computational Intelligence and Machine Learning (2026), Workshop on Explainable AI at International Joint Conference on Artificial Intelligence (XAI @ IJCAI) (2023-2024), Workshop on Cognitive Aspects of Knowledge Representation at Knowledge Representation (CAKR @ KR) (2023)

*Special recognitions for outstanding reviews: CHI'26, IUI'26, ARR May'25

DEPARTMENTAL SERVICE

PhD Interview Panel, Department of Computer Science & Statistics,
Trinity College, University of Dublin

May 2025

- Assessed and interviewed candidates for Irish Central Bank PhD Programme in AI & Data Science

Postdoctoral Network Committee, Department of Computer Science,
University of Copenhagen

Nov 2024–

- Organised career-support and networking events for postdoctoral researchers

SKILLS

Quantitative Research Methods: Designing and conducting large-scale psychological & behavioural experiments (Prolific, Alchemer, Survey Monkey, Qualtrics) and analysis in R, Python, and SPSS.

Qualitative Research Methods: Semi-structured interviews, think-aloud studies, focus groups, analysis using grounded theory & thematic analysis, and content analysis.

Software: Python (pandas, scikit-learn, nltk), R $\text{\LaTeX}$, HTML, SPSS, NVivo.

Languages: English (Native), Irish (Native), French (B1), Danish (A2)

PUBLICATIONS

* denotes joint first authorship

Conference Proceedings

1. **Greta Warren**, Jingyi Sun, Irina Shklovski, and Isabelle Augenstein. Show me the evidence: Evaluating the role of evidence and natural language explanations in AI-supported fact-checking. In *Proceedings of the 2026 ACM Conference on Fairness, Accountability, and Transparency*, FAccT '26, page 4899–4934, New York, NY, USA, 2026. Association for Computing Machinery. doi:[10.1145/3805689.3812358](https://doi.org/10.1145/3805689.3812358)
2. Jingyi Sun*, **Greta Warren***, Irina Shklovski, and Isabelle Augenstein. Explaining sources of uncertainty in automated fact-checking. In Maria Liakata, Viviane P. Moreira, Jiajun Zhang, and David Jurgens, editors, *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 45510–45534, San Diego, California, United States, July 2026. Association for Computational Linguistics. URL: <https://aclanthology.org/2026.acl-long.2110/>, doi:[10.18653/v1/2026.acl-long.2110](https://doi.org/10.18653/v1/2026.acl-long.2110) [Core A*]
3. Will Hawkins, Kai Rawal, Jonathan Rystrom, Stratis Tsirtsis, Zihao Fu, **Greta Warren**, Eoin D. Delaney, Ryan Brown, Sandra Wachter, Brent Mittelstadt, and Chris Russell. The heterogeneous safety impacts of benign multilingual fine-tuning. In *Proceedings of the 42nd International Conference on Machine Learning (to appear)*, volume 268 of *Proceedings of Machine Learning Research*. PMLR, 6–11 Jul 2026. URL: <https://openreview.net/forum?id=5M6NoV6Hs6> [Core A*]
4. Eoin Delaney, **Greta Warren**, Derek Greene, and Mark T. Keane. Explaining time series classifications counterfactually: A human-centered evaluation. *9th AAAI/ACM Conference on Artificial Intelligence, Ethics & Society (AIES'26) (to appear)*, 2026. URL: <http://dx.doi.org/10.2139/ssrn.6883018>
5. Nadav Borenstein*, **Greta Warren***, Desmond Elliott, and Isabelle Augenstein. Can community notes replace professional fact-checkers? In Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar, editors, *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 535–552, Vienna, Austria, July 2025. Association for Computational Linguistics. URL: <https://aclanthology.org/2025.acl-short.42/>, doi:[10.18653/v1/2025.acl-short.42](https://doi.org/10.18653/v1/2025.acl-short.42) [Core A*]
6. **Greta Warren**, Irina Shklovski, and Isabelle Augenstein. Show me the work: Fact-checkers' requirements for explainable automated fact-checking. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*, CHI '25, New York, NY, USA, 2025. Association for Computing Machinery. doi:[10.1145/3706598.3713277](https://doi.org/10.1145/3706598.3713277) [Core A*]

7. **Greta Warren**, Eoin Delaney, Christophe Gu  ret, and Mark T. Keane. Explaining multiple instances counterfactually: User tests of group-counterfactuals for XAI. In Juan A. Recio-Garcia, Mauricio G. Orozco-del Castillo, and Derek Bridge, editors, *Case-Based Reasoning Research and Development*, pages 206–222, Cham, 2024. Springer Nature Switzerland. doi:10.1007/978-3-031-63646-2_14
8. **Greta Warren**, Ruth M. J. Byrne, and Mark T. Keane. Categorical and continuous features in counterfactual explanations of AI systems. In *Proceedings of the 28th International Conference on Intelligent User Interfaces*, IUI ’23, pages 171–187, New York, NY, USA, 2023. Association for Computing Machinery. doi:10.1145/3581641.3584090 [Core A]
9. **Greta Warren**, Barry Smyth, and Mark T. Keane. Better counterfactuals, ones people can understand: Psychologically-plausible case-based counterfactuals using categorical features for explainable AI (XAI). In Mark T. Keane and Nirmalie Wiratunga, editors, *Case-Based Reasoning Research and Development: 30th International Conference, ICCBR 2022, Nancy, France, September 12–15, 2022, Proceedings*, pages 63–78, Berlin, Heidelberg, 2022. Springer-Verlag. doi:10.1007/978-3-031-14923-8_5

Journal Articles

1. **Greta Warren**, Ruth M. J. Byrne, and Mark T. Keane. Categorical and continuous features in counterfactual explanations of AI systems. *ACM Trans. Interact. Intell. Syst.*, 14(4), December 2024. doi:10.1145/3673907
2. Catherine N. Moran, David P. McGovern, **Greta Warren**, R  n  n    Gr  laigh, Joanne P. M. Kenney, Alan Smeaton, and Paul M. Dockree. Young and restless, old and focused: Age-differences in mind-wandering frequency and phenomenology. *Psychology and aging*, 36(2):252, 2021. doi:10.1037/pag0000526

Peer-reviewed Workshops and Symposia

1. Nadav Borenstein*, **Greta Warren***, Desmond Elliott, and Isabelle Augenstein. Can community notes replace professional fact-checkers? In *CHI’25 Workshop on News Futures*, 2025. URL: <https://arxiv.org/abs/2502.14132>, arXiv:2502.14132
2. **Greta Warren**, Mark T. Keane, Christophe Gueret, and Eoin Delaney. Explaining groups of instances counterfactually for XAI: A use case, algorithm and user study for group-counterfactuals. *IJCAI-23 Workshop on Explainable Artificial Intelligence (XAI)*, 2023. doi:10.48550/arXiv.2303.09297
3. **Greta Warren**, Mark T. Keane, and Ruth M. J. Byrne. Features of explainability: How users understand counterfactual and causal explanations for categorical and continuous features in XAI. In *IJCAI-ECAI’22 Workshop: Cognitive Aspects of Knowledge Representation*, 2022. doi:10.48550/arXiv.2204.10152
4. J  rg Cassens, Lorenz Habenicht, Julian Blohm, Rebekah Wegener, Joanna Korman, Sangeet Khemlani, Giorgio Gronchi, Ruth M. J. Byrne, **Greta Warren**, Molly S. Quinn, and Mark T. Keane. Explanation in human thinking. In *Proceedings of the 43rd Annual Meeting of the Cognitive Science Society*, 2021. URL: <https://escholarship.org/content/qt9k6291nk/qt9k6291nk.pdf>

Preprints

1. Isabelle Augenstein, Michiel Bakker, Tanmoy Chakraborty, David Corney, Emilio Ferrara, Iryna Gurevych, Scott Hale, Eduard Hovy, Heng Ji, Irene Larraz, Filippo Menczer, Preslav Nakov, Paolo Papotti, Dhruv Sahnan, **Greta Warren**, and Giovanni Zagni. Community moderation and the new epistemology of fact checking on social media, 2025. URL: <https://arxiv.org/abs/2505.20067>, arXiv:2505.20067
2. **Greta Warren**, Mark T. Keane, and Ruth M.J. Byrne. Preferences for common-cause diagnoses and common-effect predictions. 2024. URL: <https://osf.io/hvc5t>

TALKS, PANELS AND PRESENTATIONS

Invited Talks

May, 2026. Show me the evidence: Evaluating the role of evidence and natural language explanations in AI-supported fact-checking. *Workshop on LLMs and Access to Political Information, University College Dublin, Dublin, Ireland*

April, 2026. Designing Human-Centred AI for Explainable and Reliable Information Systems. *Last Fridays Talks, Pioneer Centre for AI, Copenhagen, Denmark*

December, 2025. “I’m not sure because...”: Communicating uncertainty in human-AI collaboration. *Metacognition in Generative AI Workshop, EurIPS Conference, Copenhagen, Denmark*

December, 2022. Simplicity and complexity in explanations of diagnoses and predictions. *Reasoning and Imagination Lab, Trinity College Dublin, Ireland*

November, 2022. Group counterfactual explanations for AI predictions. *Accenture Labs at The Dock, Dublin, Ireland*

Panel Discussions

December, 2025. Panel on Metacognition in humans and generative AI. Panellists: Moritz Hardt, Kai Sandbrink, Gaël Varoquaux, **Greta Warren**, Nick Yeung. *Metacognition in Generative AI Workshop, EurIPS Conference, Copenhagen, Denmark*

Oral Presentations

April, 2025. Show Me the Work: Fact-Checkers’ Requirements for Explainable Automated Fact-Checking. *CHI Conference on Human Factors in Computing Systems, Yokohama, Japan*

April, 2025. Can community notes replace professional fact-checkers? *News Futures Workshop @ CHI 2025, Yokohama, Japan*

March, 2025. Show Me the Work: Fact-Checkers’ Requirements for Explainable Automated Fact-Checking. *Nordic Pre-CHI Conference, Copenhagen, Denmark*

March, 2023. Categorical and continuous features in counterfactual explanations of AI systems. *International Conference on Intelligent User Interfaces, Warrane/Sydney, Australia*

July, 2022. Features of explainability: How users understand counterfactual and causal explanations for categorical and continuous features in XAI. *Workshop on Cognitive Aspects of Knowledge Representation at International Joint Conference on Artificial Intelligence, Vienna, Austria*

September, 2022. Counterfactual Explanations for Explainable AI (XAI). *Doctoral Consortium at International Conference on Case-Based Reasoning, Nancy, France*

June, 2021. Counterfactual Explanations in Explainable AI. *Annual Meeting of the Cognitive Science Society, Online*

Poster Presentations

December, 2025. Can community notes replace professional fact-checkers? *ELLIS Unconference, Copenhagen, Denmark*

July, 2025. Can community notes replace professional fact-checkers? *Association for Computational Linguistics (ACL), Vienna, Austria*

July, 2025. Can community notes replace professional fact-checkers? *Pre-ACL Workshop, Copenhagen, Denmark*

June, 2025. Can community notes replace professional fact-checkers? *Copenhagen NLP Symposium, Copenhagen, Denmark*

July, 2024. Explaining multiple instances counterfactually: User tests of group-counterfactuals for XAI. *International Conference on Case-Based Reasoning, Mérida, Mexico*

March, 2023. Categorical and continuous features in counterfactual explanations of AI systems. *School of Psychology Research Symposium, Trinity College Dublin, Ireland*

September, 2022. Better counterfactuals, ones people can understand: Psychologically-plausible case-based counterfactuals using categorical features for explainable AI (XAI). *International Conference on Case-Based Reasoning, Nancy, France*